

Java/Spring Boot Elasticsearch

Java/Spring Boot Elasticsearch

Estimated Time: ~ 2–3 days for Assignment A; +1 day for Assignment B (bonus)

Deliverables

1. **Git Repository** containing:
 - All source code
 - Configuration files (e.g., `docker-compose.yml`)
 - Sample data file (`sample-courses.json`)
 - A clear and step-by-step `README.md` explaining setup, indexing, and API usage
2. **Video / Screen Recording (3–5 minutes)** showing:
 - Elasticsearch running locally
 - Spring Boot application starting
 - Sample data being indexed into Elasticsearch
 - `/api/search` endpoint working with filters, pagination, and sorting
 - (Bonus) Autocomplete & fuzzy search in action
3. **Submission:**
 - Push the Git repository to GitHub/GitLab (public or private with access granted)
 - Upload or link the video/screen recording
 - Share both links via the **Google Form** – [[Click Here](#)]

Overview

Build a small Spring Boot application that:

1. Indexes a set of sample “course” documents into Elasticsearch.
2. Exposes a REST endpoint to **search** courses with multiple filters, pagination, and sorting (Assignment A).
3. (Bonus) Implements **autocomplete suggestions** and **fuzzy matching** for course titles (Assignment B).

Candidates will start from an empty project (e.g., generated via Spring Initializr) and launch Elasticsearch locally using Docker Compose.

Part 1: Elasticsearch Setup

1. **Create a `docker-compose.yml`** that spins up a single-node Elasticsearch cluster (version 7.x or 8.x).

2. Ensure the cluster is accessible on `localhost:9200` without authentication.
 3. Include instructions in your README on how to bring Elasticsearch up and verify it's running (e.g., via `curl http://localhost:9200`).
-

Part 2: Prepare Sample Data

1. Create a JSON file named `sample-courses.json` containing at least **50 course objects**. Each object must include fields for:
 - **id** (unique identifier)
 - **title** (short text)
 - **description** (longer text)
 - **category** (e.g., "Math," "Science," "Art," etc.)
 - **type** (values: `ONE_TIME`, `COURSE`, or `CLUB`)
 - **gradeRange** (e.g., "1st-3rd")
 - **minAge** and **maxAge** (numeric)
 - **price** (decimal or double)
 - **nextSessionDate** (ISO-8601 date-time string, e.g., "2025-06-10T15:00:00Z").
 2. Vary the data so you have a mix of categories, age ranges, course types, prices, and session dates spanning different weeks.
 3. Place this file under `src/main/resources` so the application can read it when bootstrapping.
-

Part 3: Assignment A (Required) – Basic Course Search with Filters

3.1. Project Initialization

- Create a **new Spring Boot project** (group/artifact of your choice).
- Include dependencies for:
 - Spring Web (REST controllers)
 - Spring Data Elasticsearch (or the official Elasticsearch client)
 - Your chosen JSON library (e.g., Jackson)
 - (Optional) Lombok or any utility libraries.

3.2. Elasticsearch Configuration

- Configure your application to connect to the local Elasticsearch instance on `localhost:9200`.
- Document in your README how to set any necessary properties (e.g., host, port) so that a reviewer can run the application without modification.

3.3. Define Course Document Structure

- Create an entity (e.g., `CourseDocument`) corresponding to the fields in `sample-courses.json`.
- Ensure each field has an appropriate Elasticsearch mapping type (e.g., `text` for searchable text, `keyword` for exact matches, `date` for session dates, `double` or `float` for price).

3.4. Bulk-Index Sample Data

- Implement a component (e.g., a service or a startup listener) that reads `sample-courses.json` at application startup and bulk-indexes all course objects into Elasticsearch's `courses` index.
- Document in your README how to trigger or verify this data ingestion.

3.5. Implement the Search Service

- Build a service layer method that executes a search query against the `courses` index, applying:
 - **Full-text search** on `title` and `description` (e.g., multi-match).
 - **Range filters** for `minAge/maxAge` and `minPrice/maxPrice`.
 - **Exact filters** for `category` and `type`.
 - **Date filter** for `nextSessionDate` (to show only courses on or after a given date).
- Implement sorting logic:
 - Default sort: ascending by `nextSessionDate` (soonest upcoming first).
 - If a `sort=priceAsc` parameter is passed, sort by `price` (low to high).
 - If `sort=priceDesc` is passed, sort by `price` (high to low).
- Add pagination support via `page` and `size` parameters.

3.6. Expose the Search Controller

- Expose a REST endpoint, for example: `GET /api/search`
- Accept query parameters:
 - `q` (search keyword)
 - `minAge`, `maxAge`
 - `category`
 - `type`
 - `minPrice`, `maxPrice`
 - `startDate` (ISO-8601)
 - `sort` (`upcoming`, `priceAsc`, `priceDesc`)
 - `page` (default 0), `size` (default 10)
- Return a JSON response containing:
 - `total` (total hits)
 - `courses` (array of matching course documents, including at least `id`, `title`, `category`, `price`, and `nextSessionDate`).

3.7. Testing & Verification

- Include in your README:
 - Exact `curl` or HTTP examples showing how to call `/api/search` with different combinations of parameters.
 - Expected behavior (e.g., correct `total` count, proper sorting, and filtering).
- (Optional, but encouraged) Provide a few basic integration tests that:
 - Spin up Elasticsearch (via Testcontainers or a local instance)
 - Index a small subset of courses
 - Verify that specific queries return the expected results.

Part 4: Assignment B (Bonus) – Autocomplete Suggestions & Fuzzy Search

If you finish Assignment A early or want extra credit, implement the following:

4.1. Autocomplete (Completion Suggester)

1. Index with a Completion Field

- Create a new Elasticsearch index (or extend the existing one) that includes a “completion” field for `title`.

- Re-index your sample data so that each course document has a “title” and a corresponding “suggest” sub-field used for autocomplete.
- 2. **Autocomplete Endpoint**
 - Expose an endpoint, for example: `GET /api/search/suggest?q={partialTitle}`
 - Query Elasticsearch using the completion suggester API to return up to 10 matching titles that start with `partialTitle`.
 - Return a JSON array of suggested course titles.
- 3. **README Examples**
 - Show in the README how to call the suggest endpoint (e.g., `q=phy`) and what responses to expect.

4.2. Fuzzy Search Enhancement

1. **Fuzzy Matching in Search**
 - Adjust your existing search logic so that, when a keyword (`q`) is provided, the `title` field can match with fuzziness (allowing small typos).
 - For instance, searching for “dinors” should still match “Dinosaurs 101.”
2. **Documentation**
 - In your README, demonstrate via examples that a fuzzy query with a typo still returns the correct document(s).

Part 5: Submission Guidelines & Upload

Git Repository

- Push all your code to a **public or private repository** (GitHub/GitLab).
- Ensure the **commit history is clear and incremental**, with separate commits for:
 - Initial setup
 - Elasticsearch indexing logic
 - Search implementation
 - Bonus features (if implemented)

README

Your `README.md` must clearly explain how to:

1. Launch Elasticsearch (`docker-compose up -d`)
2. Build and run the Spring Boot application
3. Populate the index with sample data
4. Call each endpoint with **example curl or HTTP requests**
5. (Bonus) Trigger **autocomplete and fuzzy search** queries and show sample responses

Video / Screen Recording (Mandatory)

- Record a **3–5 minute screen recording** demonstrating:
 1. Elasticsearch running locally
 2. Spring Boot application starting
 3. Sample data being indexed into Elasticsearch
 4. `/api/search` endpoint working with filters, pagination, and sorting
 5. (Bonus) Autocomplete and fuzzy search results
- Upload the **video link** (Google Drive/YouTube/unlisted link) in the form along with your Git repo.

Project Structure & Naming

- Use **clean package conventions**:
 - `config, document, repository, service, controller`
- Choose **clear class and method names**.

Tests (Optional but Recommended)

- Include **unit or integration tests** verifying core search functionality.
- If using **Testcontainers**, document how to run tests against an ephemeral Elasticsearch instance.

Evaluation Criteria

1. **Correctness** – Search results match filters and sorting rules
2. **Performance** – Efficient ES queries with filters instead of scanning all documents
3. **Code Organization** – Clear separation of concerns
4. **Documentation** – README and instructions allow setup in **≤ 30 minutes**
5. **Video Demo** – Shows application working end-to-end
6. **Bonus Points** – Working autocomplete and fuzzy search

Final Upload

- Fill out the **Google Form** - [Click Here](#)
 1. **Git repository link**
 2. **Video/screen recording link**
 3. (Optional) **Test coverage info** if included